

Data & Analytics Case Study

XKCD Analytics Pipeline



Goals and high level overview

Goal: View and measure reader insights for xkcd comics.

High level implementation:

- Pull data from API endpoint
- Transform and create reader metrics
- Store in a data warehouse or DB for analysis

Requirements:

- Portable
- Easy to run
- Reliable results


Flow Diagram and Tooling Choices

Orchestration: Apache Airflow

Database: PostgreSQL 16

Transformation: dbt Core

Infrastructure: Docker & Docker Compose



Extract and Load

Assumptions

- There is only one upload per comic
- Reader metrics are only uploaded once and not updated
- API structure is stable
- Every comic has a unique num
- Comics are published in numeric order
- API Results will always contain a D/M/Y

Extract and Load

Flow

- Calls API
- Checks if result is in DB, if not, load to comics table in PG.
- If ingested comic is not today's date, poll until the latest comic is in.
- If not, use a sensor to poll and detect when a new comic with today's date comes in.
- Once new comic is released continue with ingest and transforms.

Why this flow?

- Checking and updating the DB for comics that aren't today helps to catch if there was an unexpected or missed comic.
 - Not a full backfill but a nice addition for current scale
 - Useful for demos and locally running
- Sensor - Used for polling in airflow for newest comic
 - Includes a option to turn off completely. Useful for adhoc runs, demos, testing and the option to move to just scheduled runs if required.

Extract and Load

```
10 class XKCDClient:
11     """
12         XKCD API Client with exponential backoff.
13         """
14
15     def __init__(self, max_retries=5, base_delay=1, max_delay=60, timeout=10, base_url=None):
16         self.max_retries = max_retries
17         self.base_delay = base_delay
18         self.max_delay = max_delay
19         self.timeout = timeout
20         self.base_url = base_url or os.getenv("XKCD_BASE_URL", "https://xkcd.com")
21
22     def _request(self, url):
23         retries = 0
24
25         while retries < self.max_retries:
26             try:
27                 resp = requests.get(url, timeout=self.timeout)
28                 resp.raise_for_status()
29                 return resp.json()
30             except requests.exceptions.RequestException as e:
31                 retries += 1
32                 if retries == self.max_retries:
33                     logger.error(f"Max retries reached: {e}")
34                     raise
35
36             # exponential backoff
37             delay = min(self.max_delay, self.base_delay * (2 ** (retries - 1)))
38
39             logger.warning(
40                 f"Request failed ({e}). "
41                 f"Retrying ({retries})/{self.max_retries} in {delay}s"
42             )
43             time.sleep(delay)
44
45     def get_latest_comic(self):
46         url = f"{self.base_url}/info.0.json"
47         return self._request(url)
```

```
8     def insert_comic_if_not_exists(engine, comic_data):
9         """Insert comic into database if it doesn't already exist."""
10        Session = sessionmaker(bind=engine)
11        session = Session()
12
13        try:
14            existing = session.query(Comic).filter(Comic.num == comic_data["num"]).first()
15            if existing:
16                logger.info(f"Comic {comic_data['num']} already exists")
17                return False
18
19            comic = Comic(
20                num=comic_data["num"],
21                title=comic_data["title"],
22                transcript=comic_data.get("transcript", ""),
23               稀疏漫画数据.get("img", ""),
24               稀疏漫画数据.get("img_hd", ""),
25                day=int(comic_data["day"]),
26                month=int(comic_data["month"]),
27                year=int(comic_data["year"]),
28                published_date=date(
29                    int(comic_data["year"]),
30                    int(comic_data["month"]),
31                    int(comic_data["day"]),
32                ),
33            )
34
35            session.add(comic)
36            session.commit()
37            logger.info(f"Successfully inserted comic {comic_data['num']}")
38            return True
39
40        except Exception as e:
41            session.rollback()
42            logger.error(f"Error inserting comic {comic_data['num']}: {e}")
43            raise
44
45        finally:
46            session.close()
```

Components

/ingest/

api.py - API client

ingest_xkcd.py - main ingestion logic, passes xcom data for comic id and date for airflow

xkcd_sensor.py - sensor for polling for newest comic

/db/

connection.py - handles session starting in PG

models.py - model definition for raw comics table

operations.py - checks if comic already exists

Considerations and Improvements

Considerations


- API limitations
- Business scalability

Improvements

- Partitioning or clustering for large raw data sets to help with scaling
- Create larger backfill solution
- More unit and integration testing
- Schema validation
- Better logging and monitoring
- Connection pooling and bulk uploads

Transform - DWH (ER) Model

- Simple Star schema
- No physical constraints as we are using ELT.
- Validation is ran through DBT tests
- dim_comic uses Type 1 SDC
- comic_key is surrogate key



Transform

```
1 -- Test to ensure customer review scores are within valid range (1.0 to 10.0)
2 -- This test will FAIL if any rows are returned (indicating invalid review scores)
3 -- Customer reviews should be constrained between 1.0 (lowest) and 10.0 (highest)
4 SELECT *
5 FROM {{ ref('fact_comic_metrics') }}
6 WHERE customer_review < 1.0 OR customer_review > 10.0
7 |
```

```
1 -- Test to ensure customer review scores are within valid range (1.0 to 10.0)
2 -- This test will FAIL if any rows are returned (indicating invalid review scores)
3 -- Customer reviews should be constrained between 1.0 (lowest) and 10.0 (highest)
4 SELECT *
5 FROM {{ ref('fact_comic_metrics') }} f
6 LEFT JOIN {{ ref('dim_comic') }} d
7 | ON f.comic_key = d.comic_key
8 WHERE d.comic_key IS NULL
```

```
1 {{ config(
2   materialized = 'table'
3 )}}
4
5 SELECT
6   d.comic_key,
7   LENGTH(d.title) * 5 AS cost_eur,
8   FLOOR(random() * 10000) AS num_views,
9   ROUND((1 + (random() * 9))::numeric, 1) AS customer_review,
10  CURRENT_TIMESTAMP AS load_timestamp
11 FROM {{ ref('dim_comic') }} d
12
13 {% if is_incremental() %}
14   WHERE d.updated_at_timestamp > (SELECT MAX(load_timestamp) FROM {{ this }})
15 {% endif %}
```

dbt workflow:

- Staging layer (stg_comics view) - clean, renamed, type-cast data.
- Dimension table (dim_comic) - static metadata about comics.
- Fact table (fact_comic_metrics) - metrics: cost_eur, views, customer_review.

Business logic implemented:

- Cost: length(title) * 5 EUR.
- Views: random() * 10000.
- Customer reviews: random 1.0–10.0.

Data Quality Checks:

- dbt tests: not_null, unique, relationships, accepted_values, expression_is_true.

Considerations and Improvements


Considerations

- Business requirements

Improvements

- Randomised metrics
 - Randomised metrics may fail certain dbt tests if misconfigured.
 - Currently not idempotent (okay for this brief but not if business requirements change)
- Data quality models (e.g. comics with missing images, comics with invalid dates)
 - Helps monitor and measure “bad” data
- Performance boosting
 - Indexing for common queries
- More complex business metrics or additional dimensions (authors, tags).

Orchestration and Infrastructure



Workflow

- Airflow DAG that triggers both ingestion and dbt transformations.
- DBT tests run automatically at the end of the DAG.
- All ran within a dockerfile along with DBT and PG

Why is it beneficial

- Replicates production environment
- Reduces manual work and custom orchestration code (e.g. restarts, polling, alerting)
- Automates all processes
- Portability

Challenges


Ingestion

- Finding a clean way to balance ad-hoc runs and scheduled runs with polling

Infrastructure

- Setting up airflow/dbt/pg docker
 - Issues with docker-compose, had to break down Airflow deployment a lot.
 - Errors with user creation and unhealthy schedulers

End Output



xkcd <xkcd> Script X

```
SELECT
    d.comic_id,
    d.title,
    d.published_date,
    f.cost_eur,
    f.num_views,
    f.customer_review
FROM dim_comic d
LEFT JOIN fact_comic_metrics f USING (comic_key);
```

dim_comic(+) 1 ×

	comic_id	title	published_date	cost_eur	num_views	customer_review
1	3,171	Geologic Core Sample	2025-11-21	100	5,425	6.5

The screenshot shows a database interface with a dark theme. At the top, there's a tab bar with "xkcd" and "<xkcd> Script X". Below the tabs is a sidebar with various icons and a "AI" button. The main area contains a SQL query window with the code provided above. Below the query is a table titled "dim_comic(+) 1 ×". The table has columns: comic_id, title, published_date, cost_eur, num_views, and customer_review. A single row is displayed with values: comic_id 3,171, title Geologic Core Sample, published_date 2025-11-21, cost_eur 100, num_views 5,425, and customer_review 6.5.

Productionisation

Further Production environment considerations:

- CI/CD
- Monitoring / alerting on DAG failures.
- Tooling choices (Use managed services)
- Container orchestration
- Data governance
- Access control (secrets manager, Airflow RBAC)
- Caching for API calls
- IAC

Thank you!